

Upload the Dataset

```
from google.colab import files  
uploaded = files.upload()
```

#Load the Dataset

```
import pandas as pd  
  
df = pd.read_csv('sentimentdataset.csv')  
df.columns = df.columns.str.strip()  
df['Sentiment'] = df['Sentiment'].str.strip()  
df.head()
```

#Data Exploration

```
df.info()  
df.describe(include='all')  
df['Sentiment'].value_counts()
```

Check for Missing Values and Duplicates

```
print("Missing values:\n", df.isnull().sum())  
print("Duplicates:", df.duplicated().sum())  
  
# Drop duplicates if needed  
df = df.drop_duplicates()
```

#Visualize a Few Features

```
import seaborn as sns  
import matplotlib.pyplot as plt  
  
sns.countplot(x='Sentiment', data=df)  
  
plt.title("Sentiment Distribution")  
plt.show()
```

Identify Target and Features

```
X = df['Text'] # Input feature
```

```
y = df['Sentiment']
```

#Convert Categorical Columns to Numerical

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
y_encoded = le.fit_transform(y) # Positive, Negative, Neutral -> 2, 0, 1 (for example)
```

#One-Hot Encoding

```
# Optional: One-hot encode sentiment labels
```

```
y_onehot = pd.get_dummies(df['Sentiment'])
```

Feature Scaling (Text Vectorization using TF-IDF)

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
tfidf = TfidfVectorizer(stop_words='english', max_df=0.7)
```

```
X_tfidf = tfidf.fit_transform(X)
```

Train-Test Split

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_tfidf, y_encoded, test_size=0.2,  
random_state=42)
```

Model Building

```
from sklearn.naive_bayes import MultinomialNB
```

```
model = MultinomialNB()
```

```
model.fit(X_train, y_train)
```

Evaluation

```
from sklearn.metrics import classification_report, accuracy_score

# Get unique classes from y_test and y_pred
unique_classes = sorted(list(set(y_test) | set(y_pred)))

# Print accuracy
print("Accuracy:", accuracy_score(y_test, y_pred))

# Generate and print the classification report
print(classification_report(y_test, y_pred, target_names=[str(c) for c in unique_classes]))

# Convert target_names to strings to avoid warning
```

#Make Predictions from New Input

```
def predict_sentiment(text):
    vector = tfidf.transform([text])
    pred = model.predict(vector)[0]
    return le.inverse_transform([pred])[0]

predict_sentiment("I love this new update!")

np.int64(158)
```

Predict the Final Grade

```
final_accuracy = accuracy_score(y_test, y_pred)
print(f"Model Grade: {final_accuracy*100:.2f}%")
```

#Deployment — Building an Interactive App

```
pip install gradio pandas scikit-learn
```

Create a Prediction Function

```
def sentiment_app(text):
    vector = tfidf.transform([text])
    prediction = model.predict(vector)[0]
```

Create the Gradio Interface

```
!pip install gradio

import gradio as gr

interface = gr.Interface(
    fn=sentiment_app,
    inputs="text",
    outputs="text",
    title="🗯️ Social Media Sentiment Analyzer",
    description="Enter a comment or post and get the predicted sentiment."
)

interface.launch(share=True)
```